

Class 14: RNASeq mini Project

Le, Sarah (PID: A18518276)

Table of contents

Background	1
Data Import	1
Remove zero count genes	4
DESeq Analysis	4
Data Visualization	6
Add annotation	8
Pathway Analysis	10
KEGG Pathways	10
GO terms	25
Reactome	26
Save our results	27

Background

Here we work through a complete RNASeq analysis project. The input data comes from a knock-down experiment of a HOX gene.

Data Import

Reading the counts and metadata CSV files

```
countData <- read.csv("GSE37704_featurecounts.csv",row.names = 1)
metadata <- read.csv("GSE37704_metadata.csv")
```

```
head(countData)
```

	length	SRR493366	SRR493367	SRR493368	SRR493369	SRR493370
ENSG00000186092	918	0	0	0	0	0
ENSG00000279928	718	0	0	0	0	0
ENSG00000279457	1982	23	28	29	29	28
ENSG00000278566	939	0	0	0	0	0
ENSG00000273547	939	0	0	0	0	0
ENSG00000187634	3214	124	123	205	207	212
		SRR493371				
ENSG00000186092		0				
ENSG00000279928		0				
ENSG00000279457		46				
ENSG00000278566		0				
ENSG00000273547		0				
ENSG00000187634		258				

```
head(metadata)
```

	id	condition
1	SRR493366	control_sirna
2	SRR493367	control_sirna
3	SRR493368	control_sirna
4	SRR493369	hoxa1_kd
5	SRR493370	hoxa1_kd
6	SRR493371	hoxa1_kd

Q1. Complete the code below to filter countData to exclude genes (i.e. rows) where we have 0 read count across all samples (i.e. columns).

```
# Note we need to remove the odd first $length col
countData <- as.matrix(countData[,-1])
head(countData)
```

	SRR493366	SRR493367	SRR493368	SRR493369	SRR493370	SRR493371
ENSG00000186092	0	0	0	0	0	0
ENSG00000279928	0	0	0	0	0	0
ENSG00000279457	23	28	29	29	28	46
ENSG00000278566	0	0	0	0	0	0
ENSG00000273547	0	0	0	0	0	0
ENSG00000187634	124	123	205	207	212	258

Q2. Complete the code below to filter countData to exclude genes (i.e. rows) where we have 0 read count across all samples (i.e. columns).

```
# Filter count data where you have 0 read count across all samples.
countData = countData[rowSums(countData)>0, ]
head(countData)
```

	SRR493366	SRR493367	SRR493368	SRR493369	SRR493370	SRR493371
ENSG00000279457	23	28	29	29	28	46
ENSG00000187634	124	123	205	207	212	258
ENSG00000188976	1637	1831	2383	1226	1326	1504
ENSG00000187961	120	153	180	236	255	357
ENSG00000187583	24	48	65	44	48	64
ENSG00000187642	4	9	16	14	16	16

Some book-keeping is required as there looks to be a mis-match between metadata rows and counts columns.

```
ncol(countData)
```

```
[1] 6
```

```
nrow(metadata)
```

```
[1] 6
```

We need to remove the first column for count.

```
cleancounts <- countData[,-1]
head(cleancounts)
```

	SRR493367	SRR493368	SRR493369	SRR493370	SRR493371
ENSG00000279457	28	29	29	28	46
ENSG00000187634	123	205	207	212	258
ENSG00000188976	1831	2383	1226	1326	1504
ENSG00000187961	153	180	236	255	357
ENSG00000187583	48	65	44	48	64
ENSG00000187642	9	16	14	16	16

```
all(colnames(cleancounts) == metadata$id)
```

Warning in colnames(cleancounts) == metadata\$id: longer object length is not a multiple of shorter object length

```
[1] FALSE
```

Remove zero count genes

There are lots of genes with zero counts. We can remove these from further analysis.

```
head(cleancounts)
```

	SRR493367	SRR493368	SRR493369	SRR493370	SRR493371
ENSG00000279457	28	29	29	28	46
ENSG00000187634	123	205	207	212	258
ENSG00000188976	1831	2383	1226	1326	1504
ENSG00000187961	153	180	236	255	357
ENSG00000187583	48	65	44	48	64
ENSG00000187642	9	16	14	16	16

```
to.keep inds <- rowSums(cleancounts) >0  
nonzero_counts <- cleancounts[to.keep inds,]
```

DESeq Analysis

Load the package.

```
library(DESeq2)
```

Setup DESeq.

```
dds = DESeqDataSetFromMatrix(countData=countData,  
                             colData=metadata,  
                             design=~condition)
```

Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in
design formula are characters, converting to factors

Run DESeq.

```
dds <- DESeq(dds)
```

estimating size factors

estimating dispersions

```
gene-wise dispersion estimates
```

```
mean-dispersion relationship
```

```
final dispersion estimates
```

```
fitting model and testing
```

```
dds
```

```
class: DESeqDataSet
dim: 15975 6
metadata(1): version
assays(4): counts mu H cooks
rownames(15975): ENSG00000279457 ENSG00000187634 ...
ENSG00000276345
ENSG00000271254
rowData names(22): baseMean baseVar ... deviance maxCooks
colnames(6): SRR493366 SRR493367 ... SRR493370 SRR493371
colData names(3): id condition sizeFactor
```

Get results.

```
res = results(dds, contrast=c("condition", "hoxa1_kd", "control_sirna"))
```

Q3. Call the summary() function on your results to get a sense of how many genes are up or down-regulated at the default 0.1 p-value cutoff.

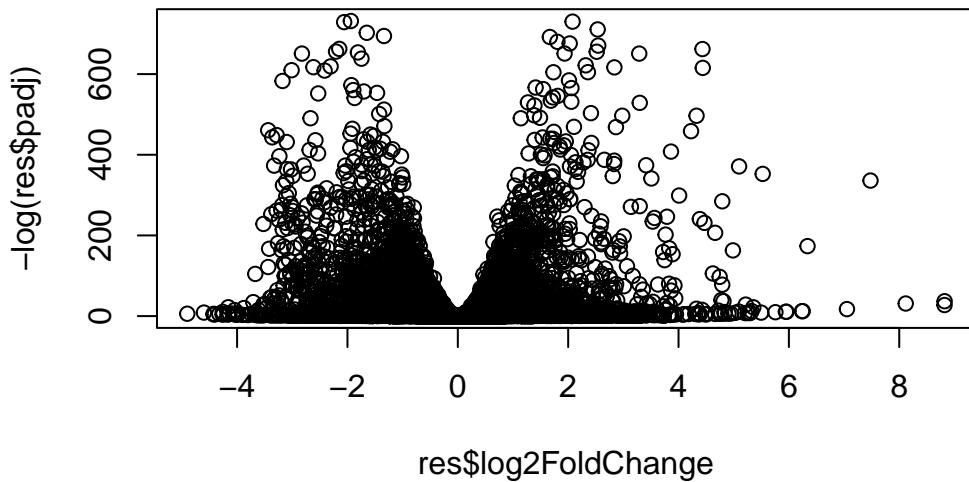
```
summary(res)
```

```
out of 15975 with nonzero total read count
adjusted p-value < 0.1
LFC > 0 (up)      : 4349, 27%
LFC < 0 (down)    : 4396, 28%
outliers [1]       : 0, 0%
low counts [2]     : 1237, 7.7%
(mean count < 0)
[1] see 'cooksCutoff' argument of ?results
[2] see 'independentFiltering' argument of ?results
```

Data Visualization

Volcano plot.

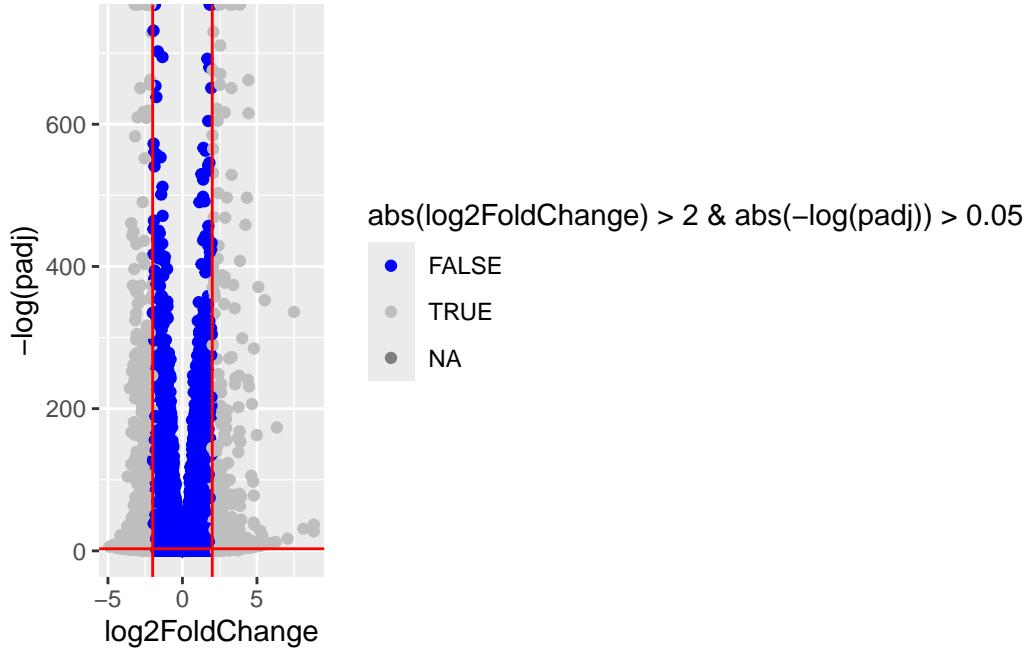
```
plot( res$log2FoldChange, -log(res$padj) )
```



```
library(ggplot2)

ggplot(res) +
  aes(log2FoldChange, -log(padj),
      color = abs(log2FoldChange) > 2 & abs(-log(padj)) > 0.05) +
  geom_point() +
  scale_color_manual(values = c("blue", "grey")) +
  geom_vline(xintercept = c (-2,2), col="red") +
  geom_hline(yintercept = -log(0.05), col="red")
```

Warning: Removed 1237 rows containing missing values or values outside the scale range
(`geom_point()`).



Add threshold lines for fold-change and p-value and color our subset of genes that make these threshold cut-offs in the plot.

Q4. Improve this plot by completing the below code, which adds color and axis labels

```

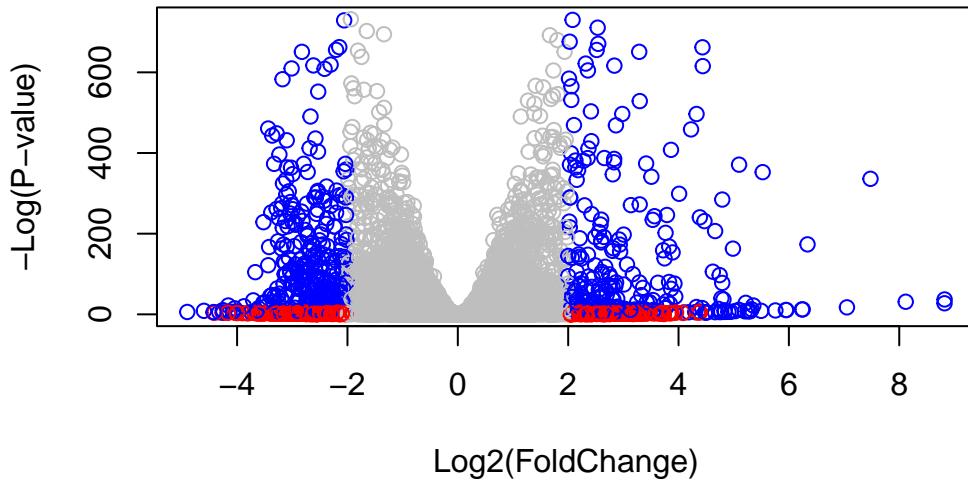
# Make a color vector for all genes
mycols <- rep("gray", nrow(res) )

# Color red the genes with absolute fold change above 2
mycols[ abs(res$log2FoldChange) > 2 ] <- "red"

# Color blue those with adjusted p-value less than 0.01
# and absolute fold change more than 2
inds <- (res$padj < 0.01) & (abs(res$log2FoldChange) > 2 )
mycols[ inds ] <- "blue"

plot( res$log2FoldChange, -log(res$padj), col= mycols, xlab="Log2(FoldChange)", ylab="-Log(Padj)" )

```



Add annotation

Add gene symbols and entrez ids.

Q5. Use the mapIDs() function multiple times to add SYMBOL, ENTREZID and GENENAME annotation to our results by completing the code below.

```
library("AnnotationDbi")
library("org.Hs.eg.db")
```

```
res$symbol <- mapIDs(x=org.Hs.eg.db,
                      keys = row.names(res),
                      keytype = "ENSEMBL",
                      column = "SYMBOL",
                      multiVals="first")
```

```
'select()' returned 1:many mapping between keys and columns
```

```

res$entrez <- mapIds(x=org.Hs.eg.db,
                      keys = row.names(res),
                      keytype = "ENSEMBL",
                      column = "ENTREZID",
                      multiVals="first")

```

'select()' returned 1:many mapping between keys and columns

```

res$name = mapIds(org.Hs.eg.db,
                  keys=row.names(res),
                  keytype= "ENSEMBL",
                  column= "GENENAME",
                  multiVals="first")

```

'select()' returned 1:many mapping between keys and columns

```
head(res, 10)
```

```

log2 fold change (MLE): condition hoxa1_kd vs control_sirna
Wald test p-value: condition hoxa1 kd vs control sirna
DataFrame with 10 rows and 9 columns
      baseMean log2FoldChange      lfcSE      stat      pvalue
      <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG00000279457    29.913579     0.1792571  0.3248216  0.551863 5.81042e-01
ENSG00000187634   183.229650     0.4264571  0.1402658  3.040350 2.36304e-03
ENSG00000188976  1651.188076    -0.6927205  0.0548465 -12.630158 1.43990e-36
ENSG00000187961   209.637938     0.7297556  0.1318599  5.534326 3.12428e-08
ENSG00000187583   47.255123     0.0405765  0.2718928  0.149237 8.81366e-01
ENSG00000187642   11.979750     0.5428105  0.5215598  1.040744 2.97994e-01
ENSG00000188290   108.922128    2.0570638  0.1969053  10.446970 1.51282e-25
ENSG00000187608   350.716868    0.2573837  0.1027266  2.505522 1.22271e-02
ENSG00000188157   9128.439422   0.3899088  0.0467163  8.346304 7.04321e-17
ENSG00000237330    0.158192     0.7859552  4.0804729  0.192614 8.47261e-01
      padj      symbol      entrez          name
      <numeric> <character> <character> <character>
ENSG00000279457 6.86555e-01        NA        NA          NA
ENSG00000187634 5.15718e-03      SAMD11    148398 sterile alpha motif ..
ENSG00000188976 1.76549e-35       NOC2L     26155 NOC2 like nucleolar ..
ENSG00000187961 1.13413e-07      KLHL17    339451 kelch like family me..
ENSG00000187583 9.19031e-01      PLEKHN1   84069 pleckstrin homology ..

```

ENSG00000187642	4.03379e-01	PERM1	84808 PPARGC1 and ESRR ind..
ENSG00000188290	1.30538e-24	HES4	57801 hes family bHLH tran..
ENSG00000187608	2.37452e-02	ISG15	9636 ISG15 ubiquitin like..
ENSG00000188157	4.21963e-16	AGRN	375790 agrin
ENSG00000237330	NA	RNF223	401934 ring finger protein ..

Q6. Finally for this section let's reorder these results by adjusted p-value and save them to a CSV file in your current project directory.

```
res = res[order(res$pvalue),]
write.csv(res, file="deseq_results.csv")
```

Pathway Analysis

KEGG Pathways

Run gage analysis with KEGG.

```
library(gage)
library(gageData)
library(pathview)

data(kegg.sets.hs)
data(sigmet.idx.hs)

# Focus on signaling and metabolic pathways only
kegg.sets.hs = kegg.sets.hs[sigmet.idx.hs]

# Examine the first 3 pathways
head(kegg.sets.hs, 3)

$`hsa00232 Caffeine metabolism`
[1] "10"    "1544"   "1548"   "1549"   "1553"   "7498"   "9"

$`hsa00983 Drug metabolism - other enzymes`
[1] "10"    "1066"   "10720"  "10941"  "151531" "1548"   "1549"   "1551"
[9] "1553"  "1576"   "1577"   "1806"   "1807"   "1890"   "221223" "2990"
[17] "3251"  "3614"   "3615"   "3704"   "51733"  "54490"  "54575"  "54576"
[25] "54577" "54578"  "54579"  "54600"  "54657"  "54658"  "54659"  "54963"
[33] "574537" "64816"  "7083"   "7084"   "7172"   "7363"   "7364"   "7365"
[41] "7366"  "7367"   "7371"   "7372"   "7378"   "7498"   "79799" "83549"
```

```
[49] "8824"   "8833"   "9"       "978"
$`hsa00230 Purine metabolism`
[1] "100"    "10201"  "10606"  "10621"  "10622"  "10623"  "107"    "10714"
[9] "108"    "10846"  "109"    "111"    "11128"  "11164"  "112"    "113"
[17] "114"    "115"    "122481" "122622" "124583" "132"    "158"    "159"
[25] "1633"   "171568" "1716"   "196883" "203"    "204"    "205"    "221823"
[33] "2272"   "22978"  "23649"  "246721" "25885"  "2618"   "26289"  "270"
[41] "271"    "27115"  "272"    "2766"   "2977"   "2982"   "2983"   "2984"
[49] "2986"   "2987"   "29922"  "3000"   "30833"  "30834"  "318"    "3251"
[57] "353"    "3614"   "3615"   "3704"   "377841" "471"    "4830"   "4831"
[65] "4832"   "4833"   "4860"   "4881"   "4882"   "4907"   "50484"  "50940"
[73] "51082"  "51251"  "51292"  "5136"   "5137"   "5138"   "5139"   "5140"
[81] "5141"   "5142"   "5143"   "5144"   "5145"   "5146"   "5147"   "5148"
[89] "5149"   "5150"   "5151"   "5152"   "5153"   "5158"   "5167"   "5169"
[97] "51728"  "5198"   "5236"   "5313"   "5315"   "53343"  "54107"  "5422"
[105] "5424"   "5425"   "5426"   "5427"   "5430"   "5431"   "5432"   "5433"
[113] "5434"   "5435"   "5436"   "5437"   "5438"   "5439"   "5440"   "5441"
[121] "5471"   "548644" "55276"  "5557"   "5558"   "55703"  "55811"  "55821"
[129] "5631"   "5634"   "56655"  "56953"  "56985"  "57804"  "58497"  "6240"
[137] "6241"   "64425"  "646625" "654364" "661"    "7498"   "8382"   "84172"
[145] "84265"  "84284"  "84618"  "8622"   "8654"   "87178"  "8833"   "9060"
[153] "9061"   "93034"  "953"    "9533"   "954"    "955"    "956"    "957"
[161] "9583"   "9615"
```

We need a named vector of fold-change values as input for gage.

```
foldchanges = res$log2FoldChange
names(foldchanges) = res$entrez
head(foldchanges)
```

```
1266      54855      1465      2034      2150      6659
-2.422719  3.201955 -2.313738 -1.888019  3.344508  2.392288
```

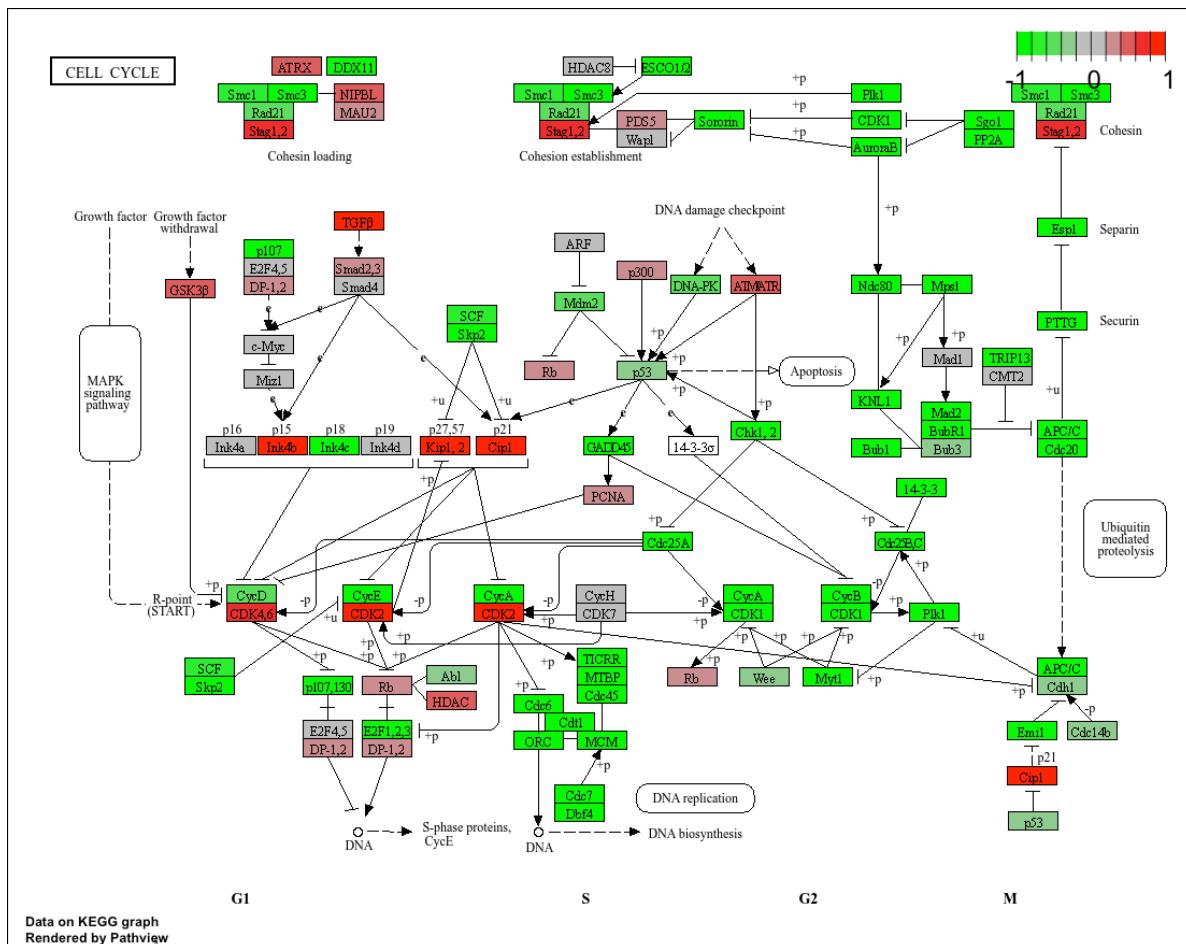
```
keggres = gage(foldchanges, gsets=kegg.sets.hs)
```

```
pathview(pathway.id = "hsa04110", gene.data = foldchanges)
```

```
'select()' returned 1:1 mapping between keys and columns
```

```
Info: Working in directory /Users/sarah/Downloads/BIMM 143/Class14
```

Info: Writing image file hsa04110.pathview.png



```
# A different PDF based output of the same data  
pathview(gene.data=foldchanges, pathway.id="hsa04110", kegg.native=FALSE)
```

'select()' returned 1:1 mapping between keys and columns

Warning: reconcile groups sharing member nodes!

```
[1,] "9"  "300"
[2,] "9"  "306"
```

Info: Working in directory /Users/sarah/Downloads/BIMM 143/Class14

```
Info: Writing image file hsa04110.pathview.pdf

## Focus on top 5 upregulated pathways here for demo purposes only
keggrespathways <- rownames(keggres$greater)[1:5]

# Extract the 8 character long IDs part of each string
keggresids = substr(keggrespathways, start=1, stop=8)
keggresids

[1] "hsa04640" "hsa04630" "hsa00140" "hsa04142" "hsa04330"

pathview(gene.data=foldchanges, pathway.id=keggresids, species="hsa")

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory /Users/sarah/Downloads/BIMM 143/Class14

Info: Writing image file hsa04640.pathview.png

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory /Users/sarah/Downloads/BIMM 143/Class14

Info: Writing image file hsa04630.pathview.png

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory /Users/sarah/Downloads/BIMM 143/Class14

Info: Writing image file hsa00140.pathview.png

'select()' returned 1:1 mapping between keys and columns

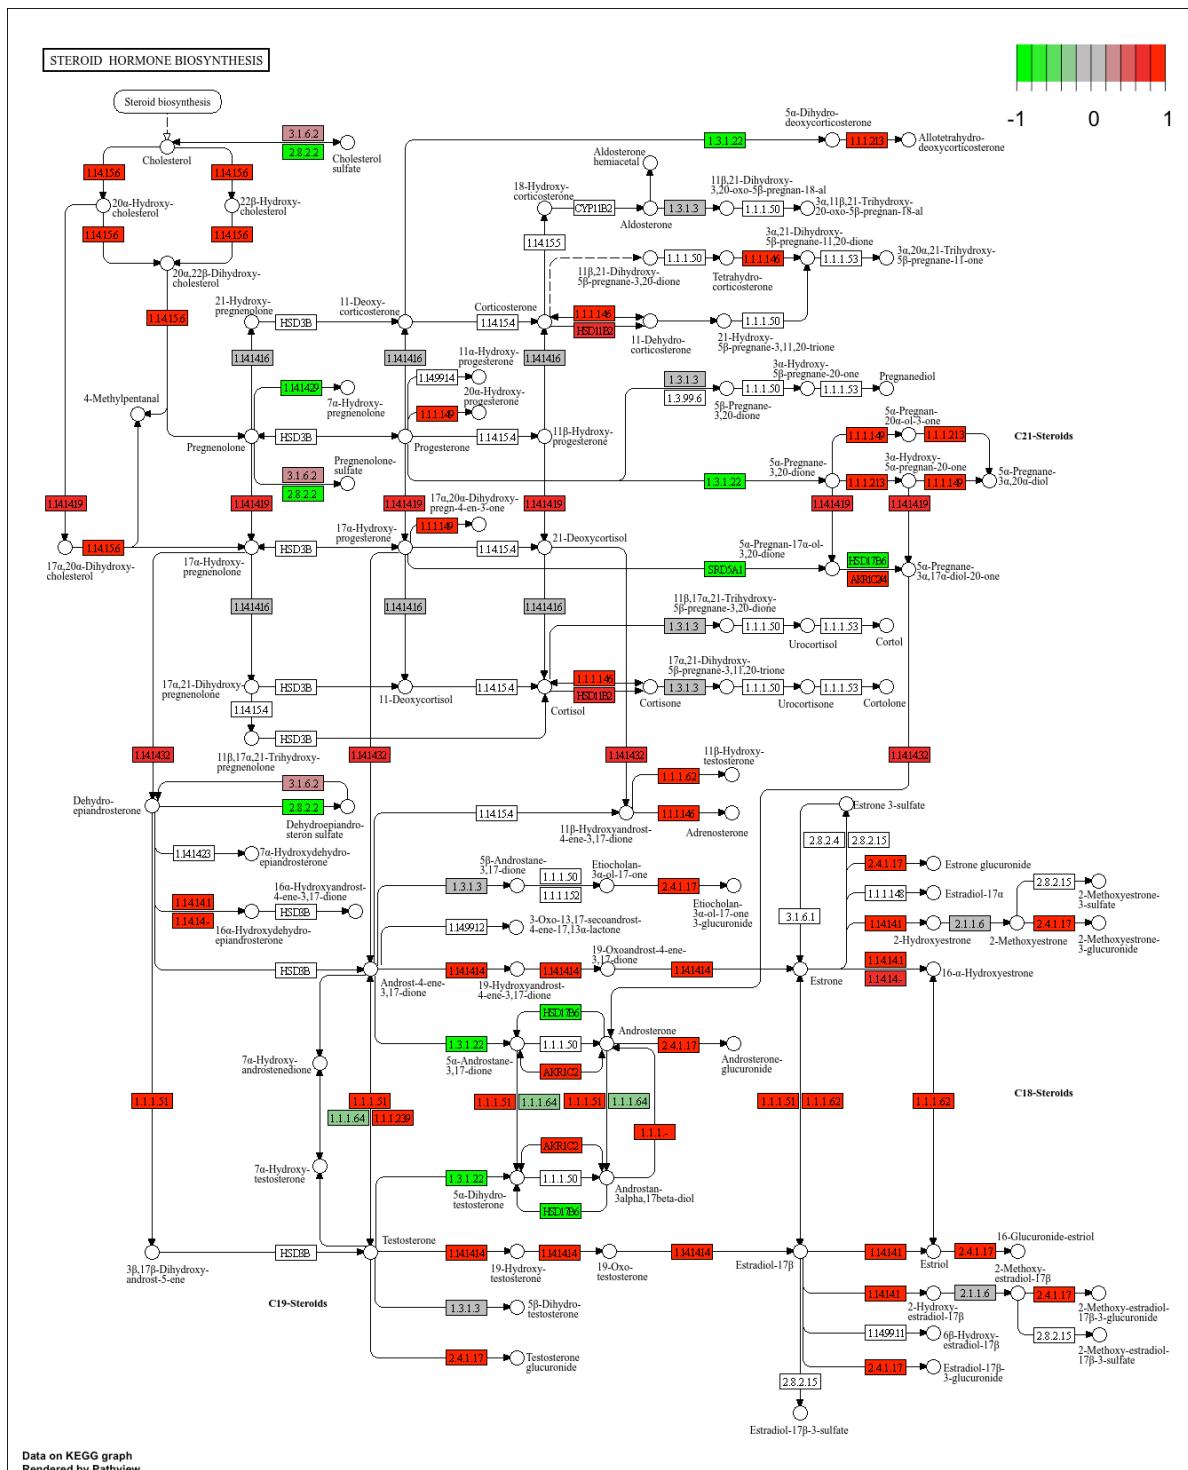
Info: Working in directory /Users/sarah/Downloads/BIMM 143/Class14

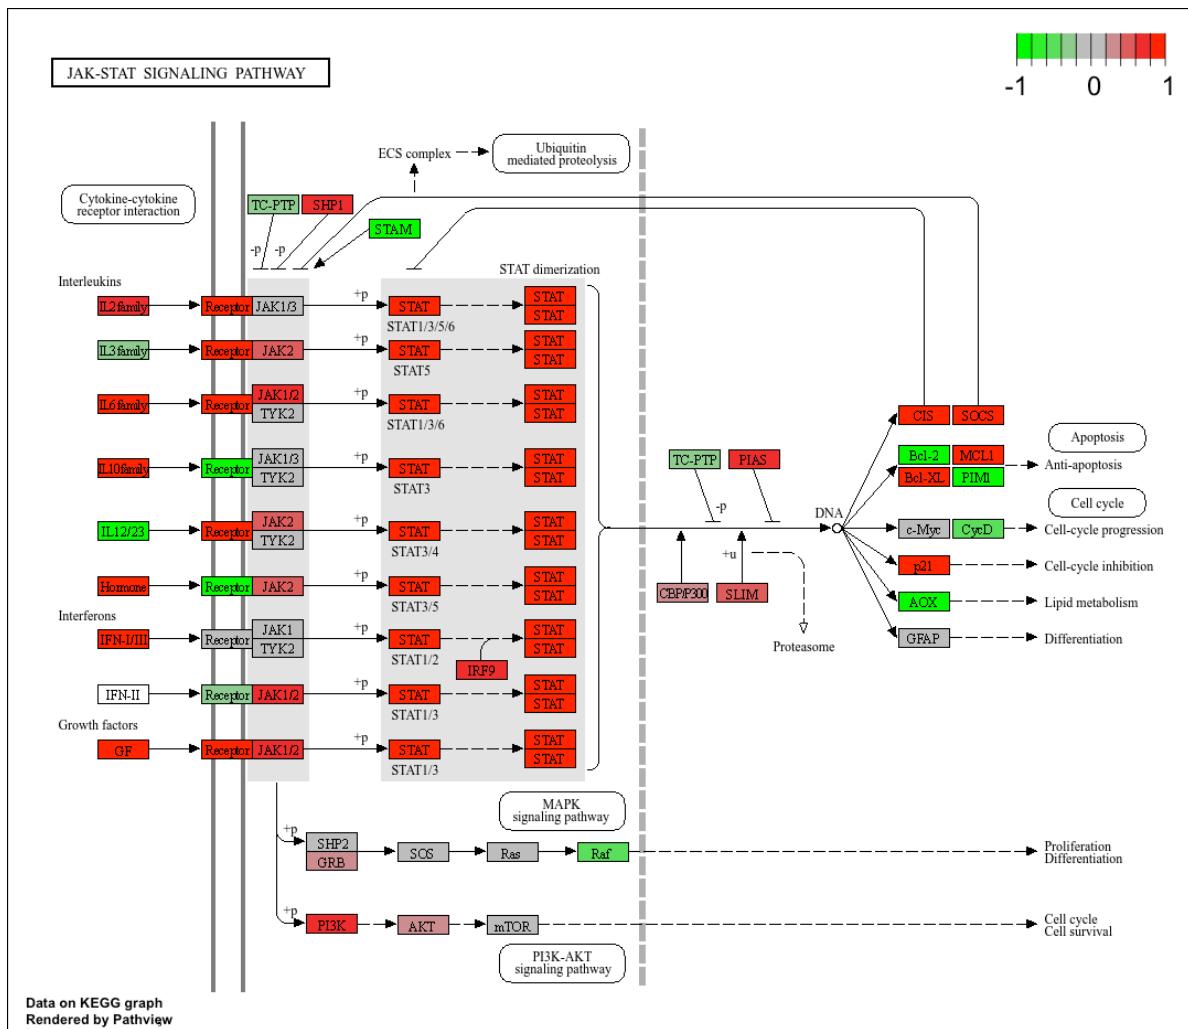
Info: Writing image file hsa04142.pathview.png
```

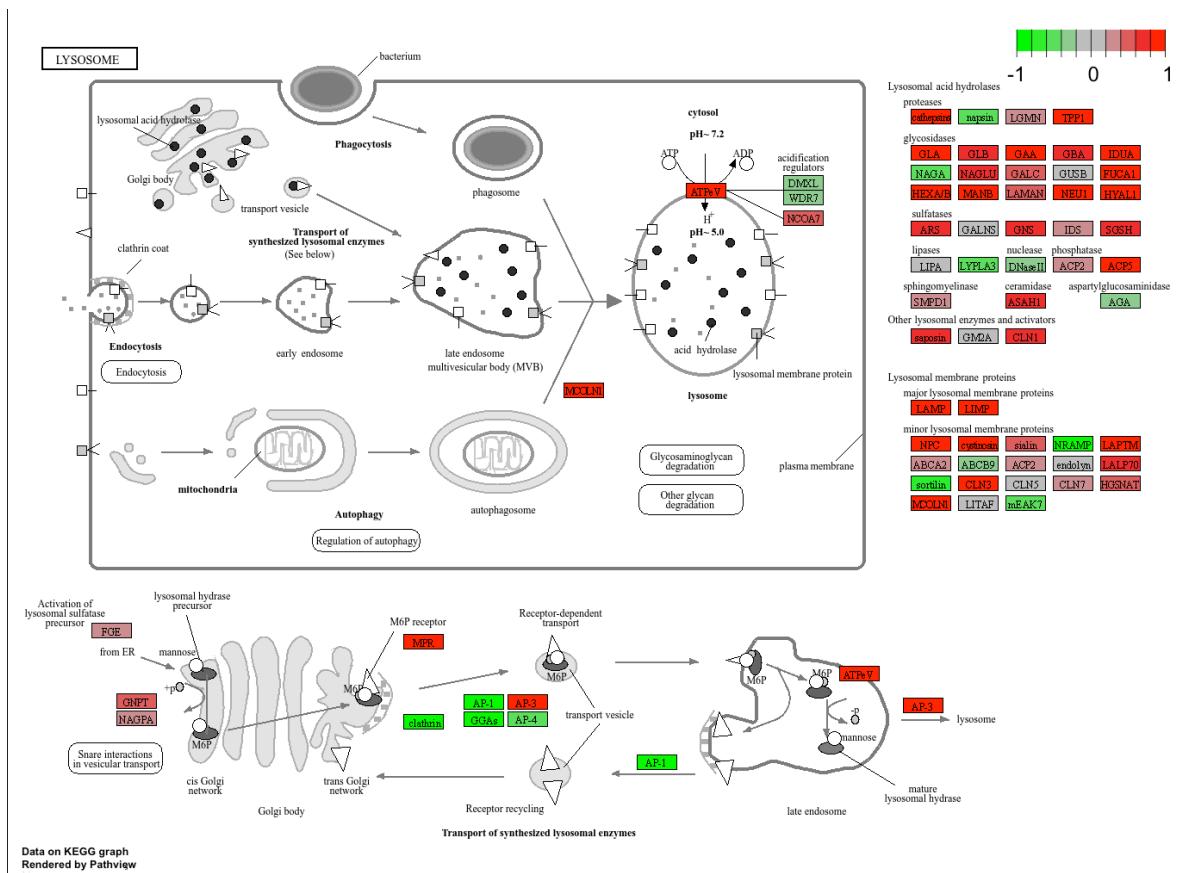
'select()' returned 1:1 mapping between keys and columns

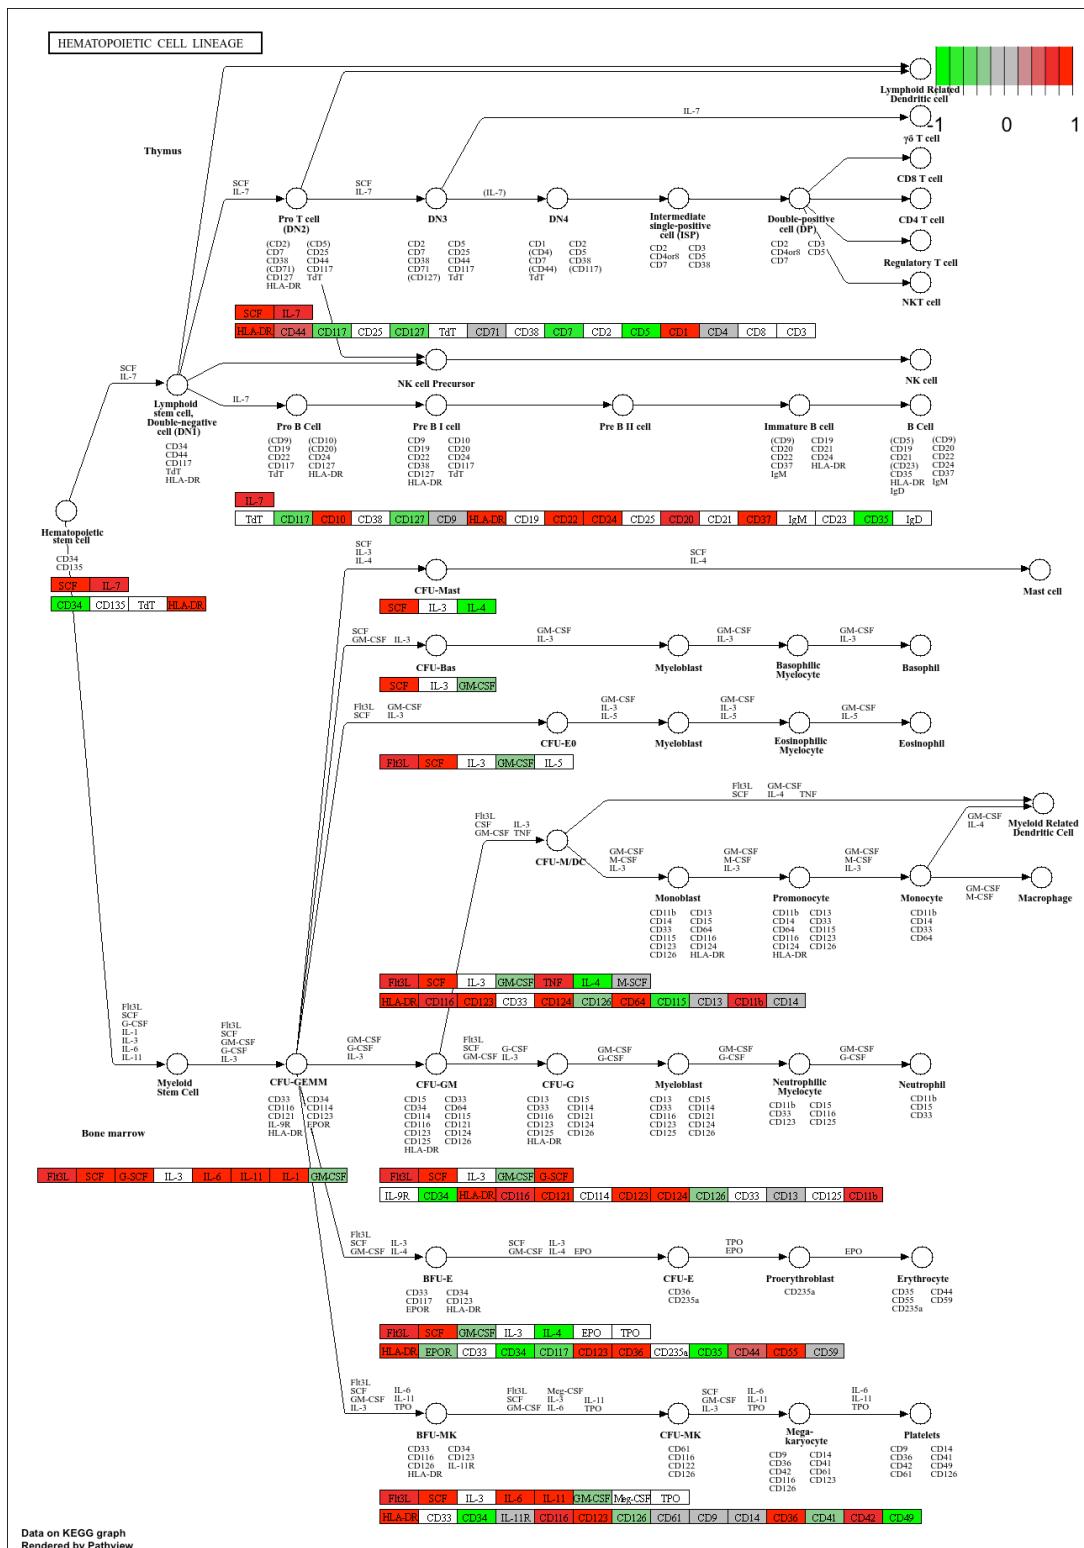
Info: Working in directory /Users/sarah/Downloads/BIMM 143/Class14

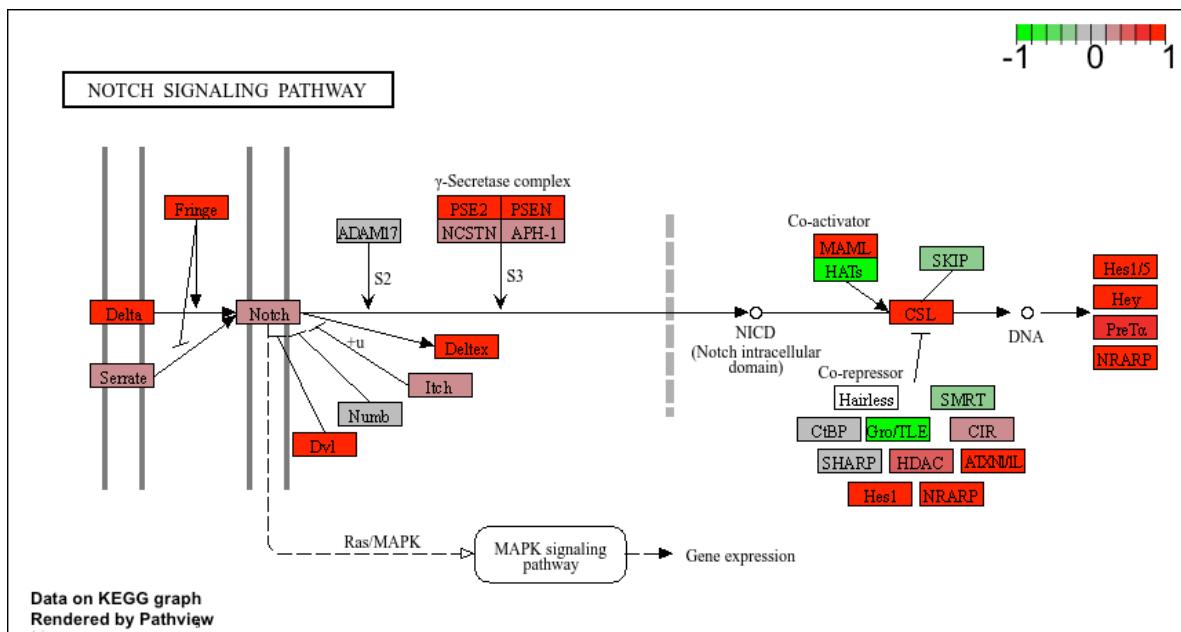
Info: Writing image file hsa04330.pathview.png











Q7. Can you do the same procedure as above to plot the pathview figures for the top 5 down-regulated pathways?

```
# Get top 5 DOWN-regulated pathways
keggrespathways_down <- rownames(keggres$less)[1:5]

# Extract only the KEGG ID part (first 8 characters)
keggresids_down <- substr(keggrespathways_down, start = 1, stop = 8)

keggresids_down
```

```
[1] "hsa04110" "hsa03030" "hsa03013" "hsa03440" "hsa04114"
```

```
pathview(gene.data = foldchanges,
          pathway.id = keggresids_down,
          species = "hsa")
```

```
'select()' returned 1:1 mapping between keys and columns
```

```
Info: Working in directory /Users/sarah/Downloads/BIMM 143/Class14
```

```
Info: Writing image file hsa04110.pathview.png
```

```
'select()' returned 1:1 mapping between keys and columns

Info: Working in directory /Users/sarah/Downloads/BIMM 143/Class14

Info: Writing image file hsa03030.pathview.png

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory /Users/sarah/Downloads/BIMM 143/Class14

Info: Writing image file hsa03013.pathview.png

'select()' returned 1:1 mapping between keys and columns

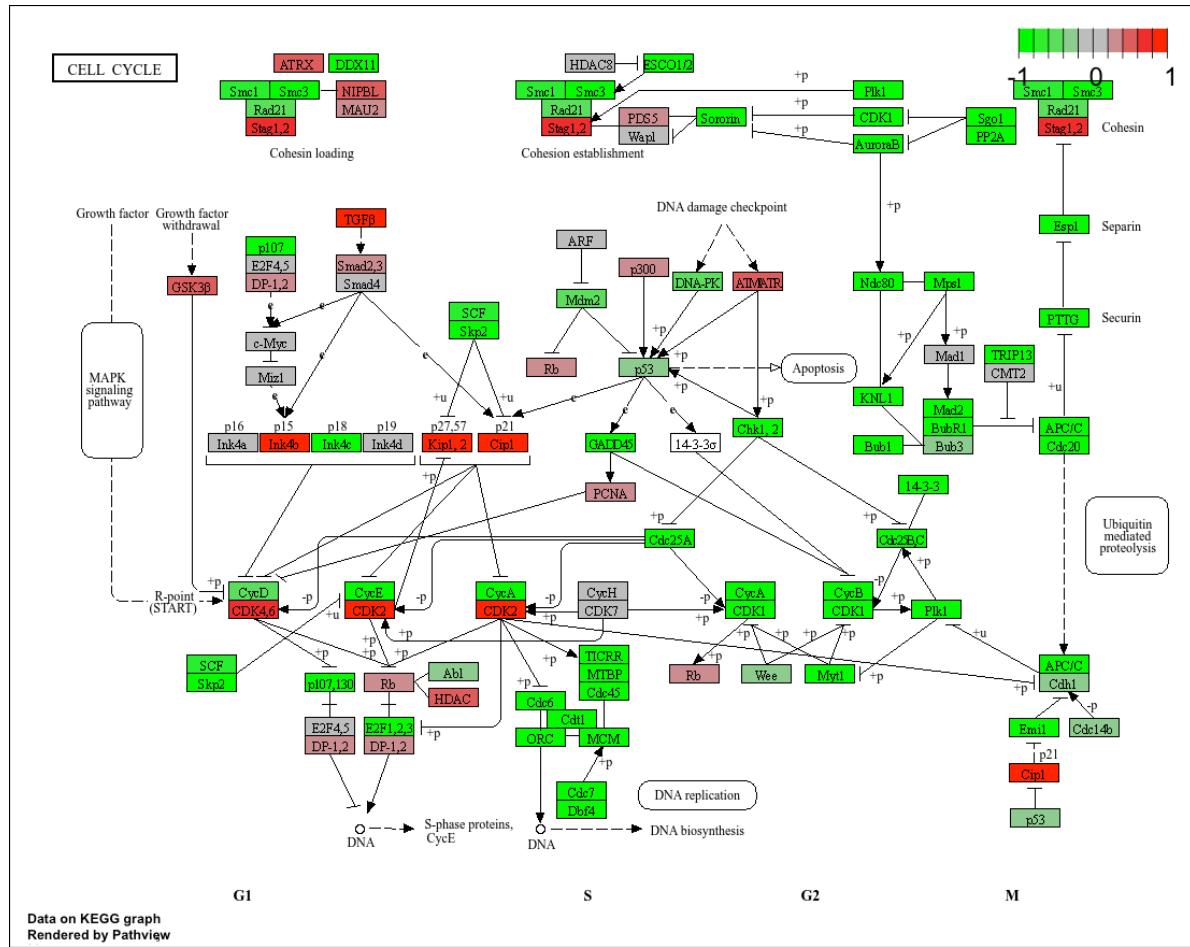
Info: Working in directory /Users/sarah/Downloads/BIMM 143/Class14

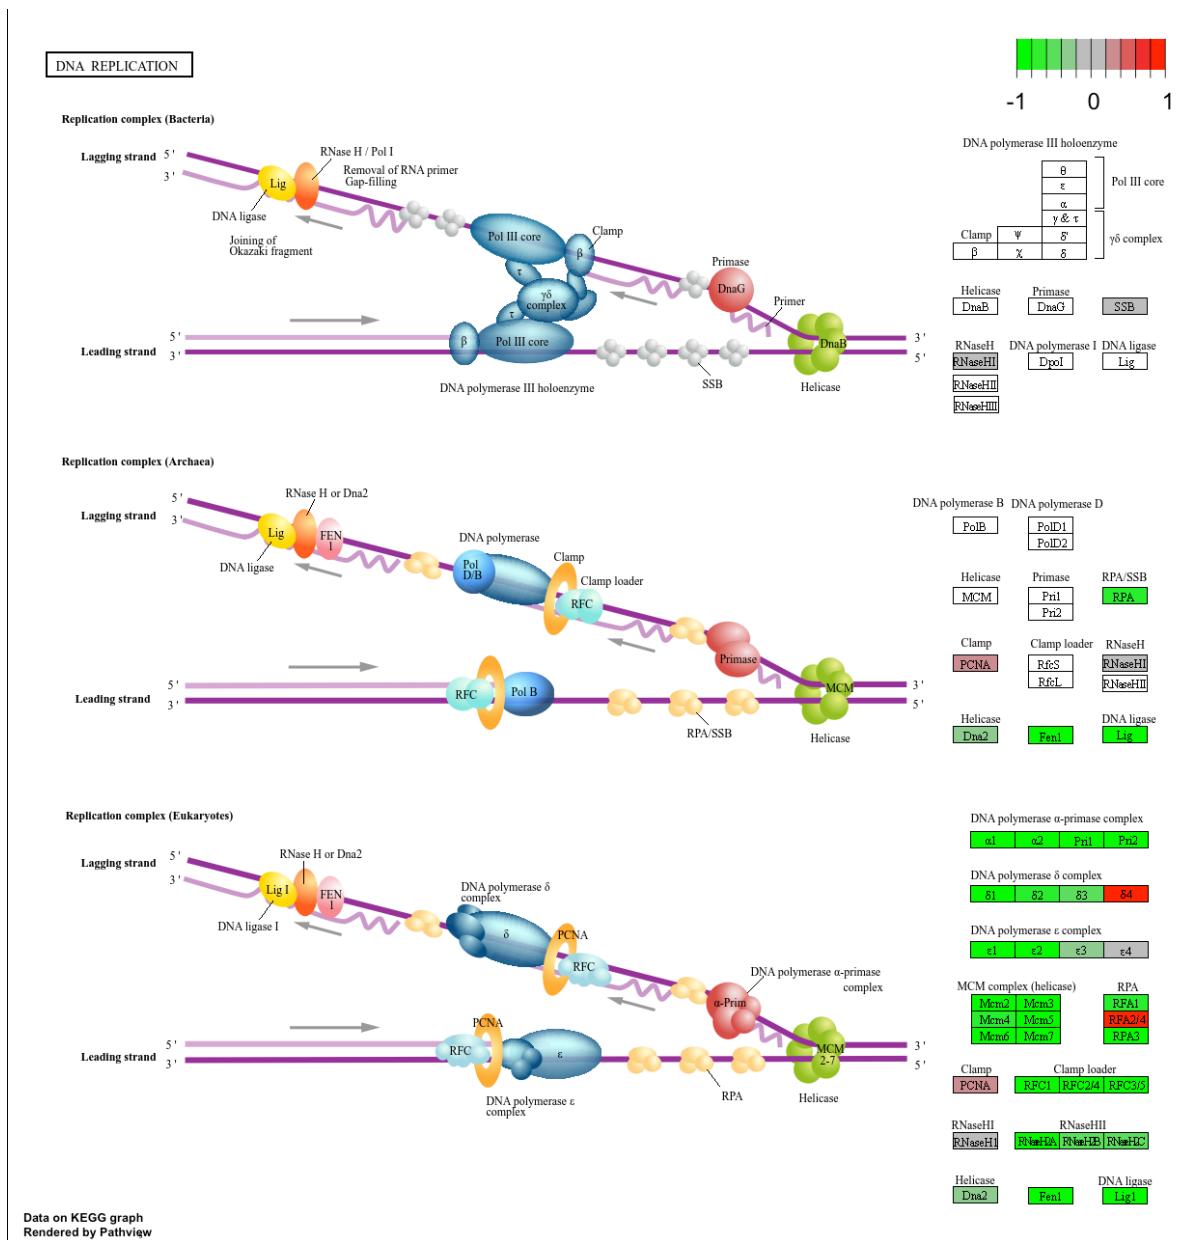
Info: Writing image file hsa03440.pathview.png

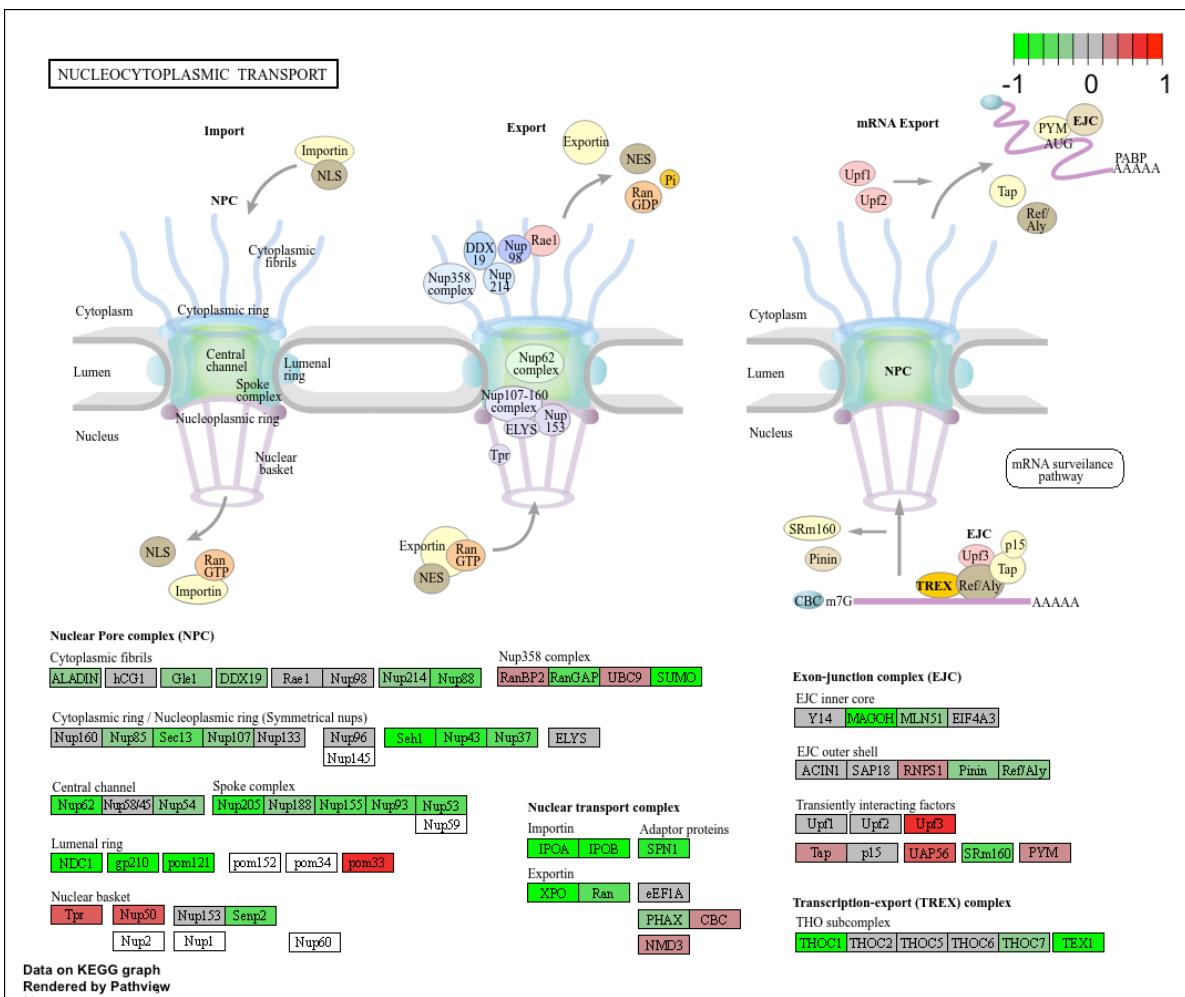
'select()' returned 1:1 mapping between keys and columns

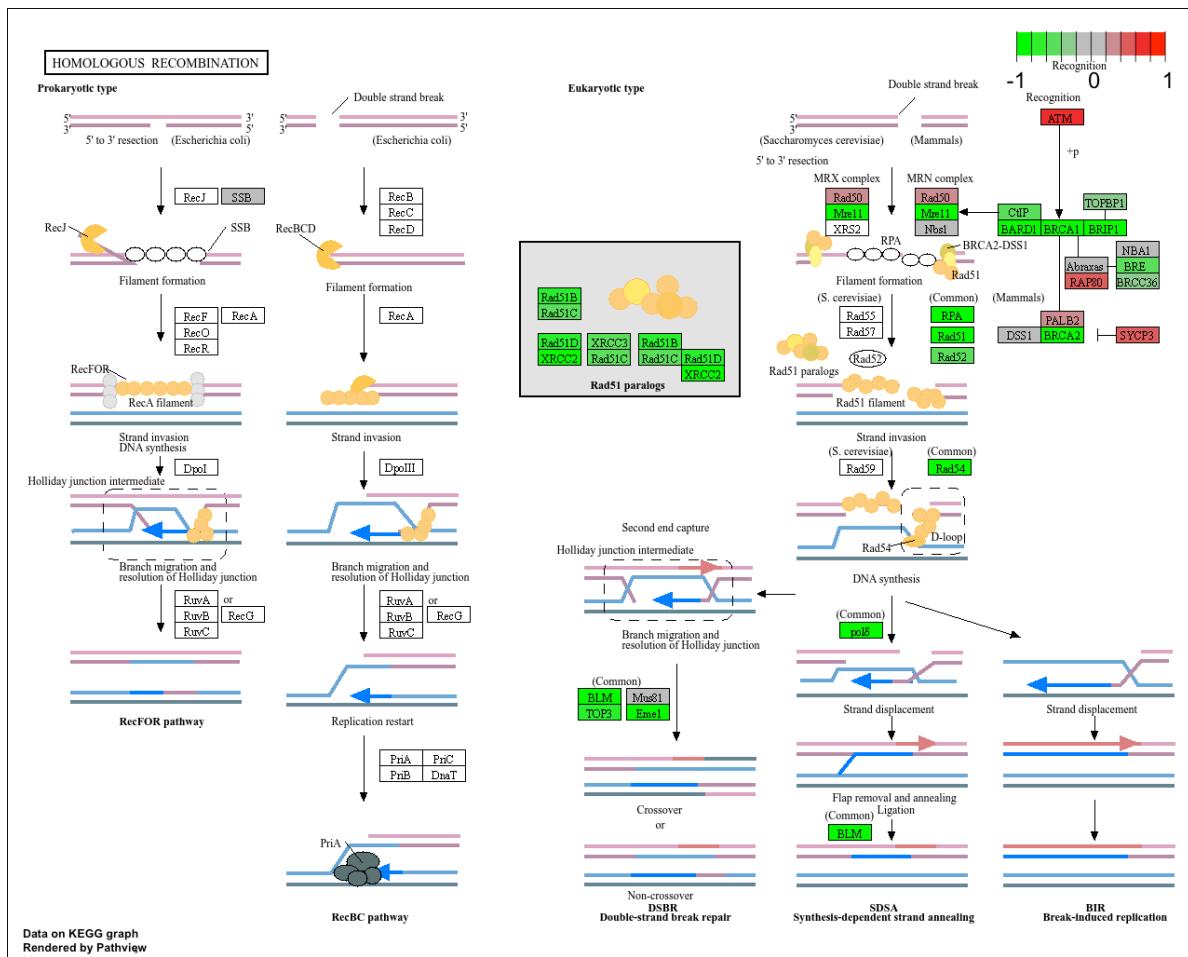
Info: Working in directory /Users/sarah/Downloads/BIMM 143/Class14

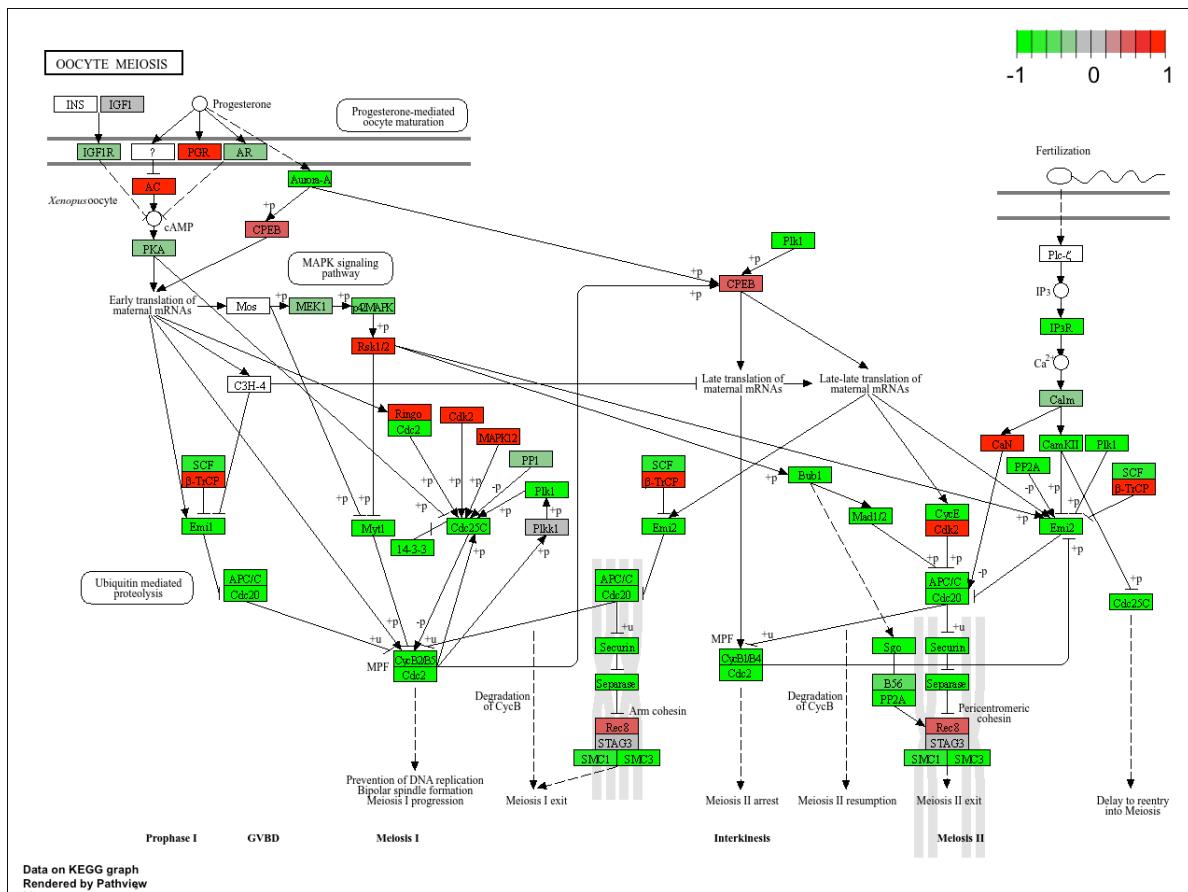
Info: Writing image file hsa04114.pathview.png
```











GO terms

Same analysis but using GO terms rather than KEGG.

```
data(go.sets.hs)
data(go.subs.hs)

# Focus on Biological Process subset of GO
gobpsets = go.sets.hs[go.subs.hs$BP]
gobpres = gage(foldchanges, gsets=gobpsets, same.dir=TRUE)
```

```
head(gobpres$less, 4)
```

	p.geomean	stat.mean	p.val
GO:0048285 organelle fission	1.536227e-15	-8.063910	1.536227e-15

GO:0000280 nuclear division	4.286961e-15	-7.939217	4.286961e-15
GO:0007067 mitosis	4.286961e-15	-7.939217	4.286961e-15
GO:0000087 M phase of mitotic cell cycle	1.169934e-14	-7.797496	1.169934e-14
	q.val	set.size	exp1
GO:0048285 organelle fission	5.841698e-12	376	1.536227e-15
GO:0000280 nuclear division	5.841698e-12	352	4.286961e-15
GO:0007067 mitosis	5.841698e-12	352	4.286961e-15
GO:0000087 M phase of mitotic cell cycle	1.195672e-11	362	1.169934e-14

Reactome

Lots of folks like the reactome web interface. You can also run this as an R function but lets look at the website first. < <https://reactome.org/>>

The website wants a text file with one gene symbol per line of the genes you want to map to pathways.

```
sig_genes <- res[res$padj <= 0.05 & !is.na(res$padj),] $symbol
head(sig_genes)
```

```
ENSG00000117519 ENSG00000183508 ENSG00000159176 ENSG00000116016 ENSG00000164251
    "CNN3"          "TENT5C"        "CSRP1"        "EPAS1"        "F2RL1"
ENSG00000124766
    "SOX4"
```

```
print(paste("Total number of significant genes:", length(sig_genes)))
```

```
[1] "Total number of significant genes: 8147"
```

And write out to a file.

```
write.table (sig_genes, file= "significant_genes.txt", row.names= FALSE, col.names=FALSE, qu
```

Q8. What pathway has the most significant “Entities p-value”? Do the most significant pathways listed match your previous KEGG results? What factors could cause differences between the two methods?

The most significant pathway is Cell Cycle (lowest entities p-value of 2.65×10^{-5}). KEGG did not have the exact same pathway but they both analyze cell cycle biology. The factors that could cause the differences between the two methods are that they use different pathway definitions, gene annotations and statistical methods. Reactome uses hierarchical pathway structures with uses only significant genes while KEGG uses broader metabolic and signaling categories with fold-change values.

Save our results

```
write.csv(res, file = "myresults.csv")
```